Yue Jin

Senior Software Engineer, Ant Group, Hangzhou, China
jinyue.derek@gmail.com — +86 15088682238 — https://jinderek.github.io/

RESEARCH INTERESTS

Machine Learning, Parallel Computing, High-performance Computing, Compiler.

EDUCATION

Zhejiang University, Hangzhou, China   Sep. 2012 — Mar. 2015
Master of Engineering in Electrical Engineering

Zhejiang University, Hangzhou, China   Sep. 2008 — Jun. 2012
Bachelor of Engineering in Electronic & Information Engineering

PUBLICATIONS

G-Sparse: compiler-driven acceleration for generalized sparse computation for graph neural networks on modern GPUs.
<u>**Y. Jin**</u>, C. Huan, H. Zhang, Y. Liu, S. L. Song, R. Zhao, Y. Zhang, C. He, W. Chen.
PACT 2023.

TEA+: a novel temporal graph random walk engine with hybrid storage architecture.
C. Huan, Y. Liu, H. Zhang, S. Song, S. Pandey, S. Chen, X. Fang, <u>**Y. Jin**</u>, B. Lepers, H. Liu, Y. Wu.
ACM TACO 2024.

GraphRPM: risk pattern mining on industrial large attributed graphs.
S. Tian, X. Zeng, Y. Hu, B. Wang, Y. Liu, <u>**Y. Jin**</u>, C. Meng, C. Hong, T. Zhang, W. Wang.
ECML PKDD 2024.

GraphGen: a distributed graph sample generation framework on industry-scale graphs.
<u>**Y. Jin**</u>, S. Tian, Y. Liu, C. Hong.
EuroSys 2024 (poster track).

GPC: compiler-based optimization for sparse computations in graph neural networks.
<u>**Y. Jin**</u>, Y. Liu.
EuroSys 2023 (poster track).

Woodpecker-DL: Accelerating Deep Neural Networks via Hardware-Aware Multifaceted Optimizations.
Y. Liu*, <u>**Y. Jin**</u>*, Y. Chen, T. Teng, H. Ou, R. Zhao, Y. Zhang.
arXiv preprint, arXiv:2008.04567, 2020.

TALKS

G-Sparse: Compiler-driven acceleration for generalized sparse computation for graph neural networks on modern GPUs.
PACT Conference 2023.

Model-based cost estimation and its application in deep learning operation optimizations.
GPU Technology Conference 2020 (GTC 2020), China.

Woodpecker-DL: An Efficient Compiler for Accelerating Deep Learning on Heterogeneous Computing Architectures.
GPU Technology Conference 2019 (GTC 2019), China.

Woodpecker Project Presentation.
Stanford Fall DAWN Retreats 2019, Quadrus, Menlo Park, CA, USA.

SELECTED PROJECTS

Compiler-based GNN Accelerator on GPU − G-Sparse   Hangzhou, China
Senior Software Engineer   Mar. 2022 — Present

- Led the development of G-Sparse, achieving a 2.4× speedup on training and inference and a 1.3× to 4.8× speedup on key operators (g-SpMM and g-SDDMM) over DGL and NVIDIA cuSparse.
- Implemented shared-memory tiling, register tiling, row balancing, and warp-shuffle techniques and extended Halide DSL into a Sparse-specific DSL.

- Published papers in PACT 2023 and contributed to open-source libraries and frameworks (Halide, TuGraph).
- Integrated into the Python package, the optimized sparse operator library can accelerate DGL and PYG training/inference with a single line of code for users.

High-performance Distributed Graph Sampling Engine - GraphGen Hangzhou, China
Senior Software Engineer Mar. 2022 — Present

- Project lead of GrpahGen, a high-performance distributed graph sampling engine, achieving 10 million nodes per second performance—20× faster than SQL-based solutions—and significantly improving sample generation for industry-scale graphs.

High-performance compiler-based Deep Learning Framework - Woodpecker Hangzhou, China
Senior Software Engineer Oct. 2028 — Mar. 2022

- Tech lead of the deep learning high-performance operator library.
- Developed a domain-specific language (DSL) compiler (based on Halide) and ML-based cost model, reducing auto-tuning time from minutes to seconds.
- Achieved a 1.5× to 10× speedup compared to NVIDIA standard libraries (cuDNN, cuBLAS, cuSPARSE) on core operations (Conv, Matmul, LayerNorm, ArgMax, ArgMin, SpMM, SDDMM, etc.).
- Achieved 1.2× to 1.7× speedup on DNN and GNN models such as ResNet-50, DeepFM, Transformer, GAT, GCN and GraphSage, based on compiler optimization and auto-fusion techniques.
- Presented work at GTC 2020 and the Stanford DAWN Retreat 2019, receiving recognition for contributions to heterogeneous computing optimizations.

INDUSTRY EXPERIENCE

Ant Group Hangzhou, China
Senior Software Engineer Oct. 2018 — Present

- Developed G-Sparse, a compiler-based GNN accelerator for GPU.
- Developed GraphGen, a high-performance distributed graph sampling engine.
- Developed Woodpecker, a high-performance compiler-based deep learning framework.

Alibaba Group Hangzhou, China
Software Engineer Apr. 2015 — Oct. 2018

- Developed JSNI, the first standardized native **I**nterface for **J**ava**S**cript and **N**ative C/C++ code interactions, widely adopted by Alibaba Group and other industry companies.
- Enhanced the Multithreaded V8 JavaScript Virtual Machine Project by optimizing its garbage collection module.

AWARDS

Most Innovative Spirit Award Hangzhou, China
Excellent Engineer: Most Innovative Spirit Award, Ant Group. 2021
Recognized for contributions to AI infrastructure and high-performance computing.

ENGLISH

TOEFL (iBT): 102 (overall score) Test date: Oct. 2024

SKILLS

- **Programming:** C/C++, CUDA, Python, JavaScript.
- **Software:** Halide, PyTorch, Tensorflow, DGL, Distributed Systems, Triton/MLIR/TVM, Node.js/V8, Git/Gerrit/Github, Linux/GNU, ARM/x86.

REFERENCES

Dr. Yao Zhang
Principal Architect at Microsoft, California, United States, **E-mail**: zhanyao@microsoft.com
Dr. Yongchao Liu
Staff Engineer at Ant Group, Hangzhou, China, **E-mail**: yongchao.ly@antgroup.com
Prof. Dekuang Su
Professor of School of Mathematical Sciences at Zhejiang University, Hangzhou, China, **E-mail**: sdk001@zju.edu.cn